

```
# Set up CUDA
#First Change runtime to GPU and run this cell
!pip install git+https://github.com/afnan47/cuda.git
%load_ext nvcc_plugin

Collecting git+https://github.com/afnan47/cuda.git
  Cloning https://github.com/afnan47/cuda.git to /tmp/pip-req-build-uu5p72
  Running command git clone --filter=blob:none --quiet https://github.com/afnan47/cuda.git /tmp/pip-req-build-uu5p72
  Resolved https://github.com/afnan47/cuda.git to commit aac710a35f52bb78ab34d2e52517237941399eff
  Preparing metadata (setup.py) ... done
Building wheels for collected packages: NVCCPlugin
  Building wheel for NVCCPlugin (setup.py) ... done
  Created wheel for NVCCPlugin: filename=NVCCPlugin-0.0.2-py3-none-any.whl size=4290 sha256=64c6f8ce1db35684c490852b0410
  Stored in directory: /tmp/pip-ephem-wheel-cache-lh84fhqd/wheels/bc/4e/e0/2d86bd15f671db32144013f1159dba09757fde36dc5
Successfully built NVCCPlugin
Installing collected packages: NVCCPlugin
Successfully installed NVCCPlugin-0.0.2
created output directory at /content/src
Out bin /content/result.out
```

```
%%writefile matrix_multiply.cu
#include <iostream>
using namespace std;

// CUDA kernel to multiply matrices
__global__ void multiply(int* A, int* B, int* C, int size) {
    int row = blockIdx.y * blockDim.y + threadIdx.y;
    int col = blockIdx.x * blockDim.x + threadIdx.x;

    if (row < size && col < size) {
        int sum = 0;
        for (int i = 0; i < size; i++) {
            sum += A[row * size + i] * B[i * size + col];
        }
        C[row * size + col] = sum;
    }
}

void initialize(int* matrix, int size) {
    for (int i = 0; i < size * size; i++) {
        matrix[i] = rand() % 10;
    }
}

void print(int* matrix, int size) {
    for (int row = 0; row < size; row++) {
        for (int col = 0; col < size; col++) {
            cout << matrix[row * size + col] << " ";
        }
        cout << '\n';
    }
    cout << '\n';
}

int main() {
    int N = 2;
    size_t matrixBytes = N * N * sizeof(int);

    int* A = new int[N * N];
    int* B = new int[N * N];
    int* C = new int[N * N];

    initialize(A, N);
    initialize(B, N);

    cout << "Matrix A:\n";
    print(A, N);

    cout << "Matrix B:\n";
    print(B, N);

    int *d_A, *d_B, *d_C;
    cudaError_t err;

    // Allocate memory on the device
    err = cudaMalloc(&d_A, matrixBytes);
    if (err != cudaSuccess) {
        cout << "CUDA malloc failed for A: " << cudaGetErrorString(err) << endl;
        return -1;
    }
}
```

```

err = cudaMalloc(&d_B, matrixBytes);
if (err != cudaSuccess) {
    cout << "CUDA malloc failed for B: " << cudaGetErrorString(err) << endl;
    return -1;
}

err = cudaMalloc(&d_C, matrixBytes);
if (err != cudaSuccess) {
    cout << "CUDA malloc failed for C: " << cudaGetErrorString(err) << endl;
    return -1;
}

// Copy data from host to device
err = cudaMemcpy(d_A, A, matrixBytes, cudaMemcpyHostToDevice);
if (err != cudaSuccess) {
    cout << "CUDA memcpy failed for A: " << cudaGetErrorString(err) << endl;
    return -1;
}

err = cudaMemcpy(d_B, B, matrixBytes, cudaMemcpyHostToDevice);
if (err != cudaSuccess) {
    cout << "CUDA memcpy failed for B: " << cudaGetErrorString(err) << endl;
    return -1;
}

// Thread and block dimensions
dim3 threads(2, 2);
dim3 blocks((N + threads.x - 1) / threads.x, (N + threads.y - 1) / threads.y);

// Launch kernel
multiply<<<blocks, threads>>>(d_A, d_B, d_C, N);

// Synchronize to make sure the kernel finishes
cudaDeviceSynchronize();

// Check for kernel launch errors
err = cudaGetLastError();
if (err != cudaSuccess) {
    cout << "CUDA kernel launch failed: " << cudaGetErrorString(err) << endl;
    return -1;
}


// Copy result back to host
err = cudaMemcpy(C, d_C, matrixBytes, cudaMemcpyDeviceToHost);
if (err != cudaSuccess) {
    cout << "CUDA memcpy failed for C: " << cudaGetErrorString(err) << endl;
    return -1;
}

// Output the result
cout << "Multiplication of Matrix A and B:\n";
print(C, N);

// Clean up
delete[] A;
delete[] B;
delete[] C;
cudaFree(d_A);
cudaFree(d_B);
cudaFree(d_C);


return 0;
}

```

 Overwriting matrix_multiply.cu

```
!nvcc -arch=sm_75 -o matrix_multiply matrix_multiply.cu
```

```
!./matrix_multiply
```

 Matrix A:
3 6
7 5

Matrix B:
3 5
6 2

Multiplication of Matrix A and B:

45 27

51 45